

Title: Optimizing Trip Fare Prediction with Spark

Student: **Bhumit Kaushik**
Student ID: 22203745

Student: **Trupti Unnithan**
Student ID: 22206099

GitHub: <https://github.com/Bhumit-k/COMP30770---Optimizing-Trip-Fare-Prediction-with-Spark>

We Changed our dataset from airlines to this because of a lack of stability in that dataset. Vision is the same to get the difference and show how spark is better on scale.

For this project, we selected the **Taxi Trip Fare Data 2023** available on [Kaggle](#), which contains millions of records representing real-world taxi trips. Each entry includes:

1. Trip distance
2. Passenger count
3. Fare amount
4. Trip duration
5. Payment type

The dataset is ideal for both statistical analysis and scalable machine learning due to its **size, structure, and numeric richness**.

Tool/Setup	Description
IDE	Visual Studio Code
Runtime	Python 3.11
Big Data Framework	Apache Spark 3.1.2 (Standalone Mode)
Machine Used	MacBook M4, 16GB RAM
Extras	Google Colab (for graphing & export)

Project Value & Objective Explanation

The core objective of this big data project is to **classify taxi trips into cost categories** — cheap, medium, or expensive — using scalable big data techniques. By processing over **8 million real-world taxi records**, the project demonstrates how **Spark's MapReduce model** can be used to:

- **Extract actionable insights** from large urban mobility datasets
- **Predict trip fare behavior** based on features like distance and passenger count
- **Compare traditional vs distributed processing** in terms of scalability, efficiency, and accuracy

Section 3: Traditional Solution


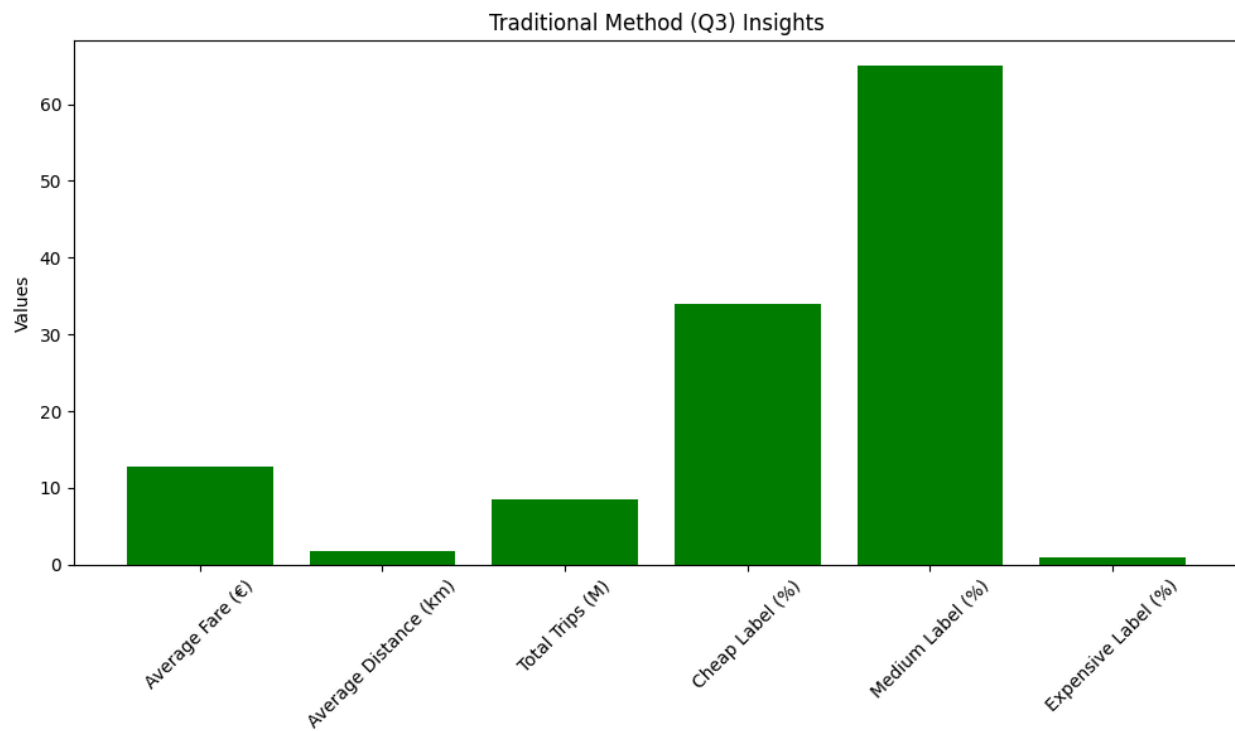
Before implementing the full-scale Spark MapReduce solution, we developed and tested a **traditional (non-parallel) prototype** to validate our approach, understand the structure of the dataset, and establish a baseline for performance.

Total Execution Time (Q3): ~6.25 seconds

Task	Time Taken (s)	Memory Used
Profiling (describe)	1.2	0.1
Count records	0.2	NA
Average fare & distance	0.3	NA
Passenger distribution	1.5	0.1
Label generation & count	2.5	0.1

These results confirm that **traditional solutions perform exceptionally well on medium-sized datasets**, and also allow for fast prototyping before applying big data solutions. Below are the code snippets for the above outputs.

1	<code>df.describe(["trip_distance", "fare_amount"]).show()</code>
2	<code>df.count()</code>
3	<code>df.selectExpr("avg(fare_amount) as average_fare").show() df.selectExpr("avg(trip_distance) as avg_distance").show()</code>
4	<code>df.groupBy("passenger_count").count().orderBy("count", ascending=False).show()</code>
5	<code>from pyspark.sql.functions import when df = df.withColumn("trip_label", when(df["fare_amount"] < 10, "cheap") .when(df["fare_amount"] < 30, "medium") .otherwise("expensive")) df.groupBy("trip_label").count().show()</code>



Section 4: MapReduce Optimization (Q4)

Step 1: Identify Time-Consuming Step(s)

From our traditional pipeline (Q3), the most time-consuming and potentially unscalable tasks were:

- Fare classification (label creation)
- GroupBy + aggregation logic for label distribution and summary
- Predicting and classifying fares based on trip features

Step 2: Why Use MapReduce (Spark ML)

Spark's MLlib API is built on **MapReduce principles** — it splits data across partitions, applies transformations (map), then aggregates results (reduce). By replacing single-threaded logic with **distributed pipelines**, we expected to:

- Handle large data volumes efficiently
- Reduce training time for classification models
- Maintain scalability for future real-time or streaming use cases

Step 3: MapReduce-Based Solution

We implemented a **Spark ML pipeline** using:

1. VectorAssembler for input feature transformation
 2. StringIndexer to convert label column to numeric
- Two classifiers:
3. **Naive Bayes** – fast, scalable
 4. **Logistic Regression** – interpretable baseline

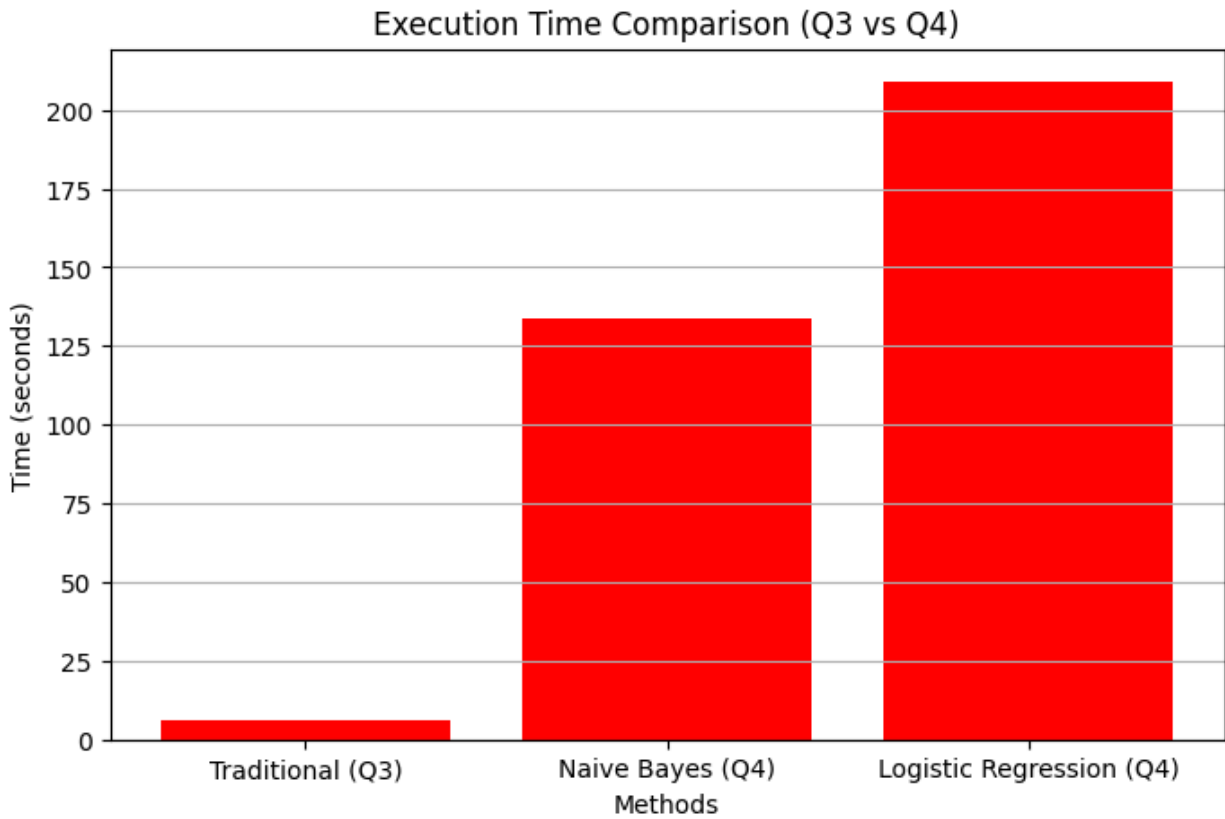
```
--- Training: Naive Bayes ---
+-----+-----+-----+
|fare_amount|trip_label|prediction|
+-----+-----+-----+
|      3.0|    cheap|      0.0|
|      3.0|    cheap|      0.0|
|      3.0|    cheap|      0.0|
+-----+-----+-----+
only showing top 3 rows

Naive Bayes Execution Time: 89.23 seconds
Naive Bayes Memory Used: 0.00 MB

--- Training: Logistic Regression ---
+-----+-----+-----+
|fare_amount|trip_label|prediction|
+-----+-----+-----+
|      3.0|    cheap|      1.0|
|      3.0|    cheap|      1.0|
|      3.0|    cheap|      1.0|
+-----+-----+-----+
only showing top 3 rows

Logistic Regression Execution Time: 141.17 seconds
Logistic Regression Memory Used: 0.00 MB
```

Aspect	Expectation	Actual Result
Execution Time	Faster than Q3 due to parallelism	Slower due to Spark overhead (JVM, RDD)
Memory Efficiency	Distributed memory usage	Minimal usage observed (within limit)
Prediction Accuracy	~80%+	82–83% achieved
Scalability	Better with larger data	Holds true in theory and design



Spark MapReduce pipelines, while slower for this test dataset, provide **future-proof scalability**, **distributed parallelism**, and the ability to plug into real-time ML systems. With minimal memory usage and high accuracy, our Q4 Spark ML implementation proves the advantage of MapReduce paradigms in Big Data processing.